

Data Mining 2025

Exercises Naive Bayes and Logistic Regression

Exercise 1: Multinomial Naive Bayes for Text Classification

You are given the following collection of recipes with corresponding kitchen:

recipeID	words in recipe	Kitchen
r1	spaghetti tomato minced meat	Italian
r2	spaghetti gorgonzola eggplant zucchini	Italian
r3	spaghetti olive oil garlic	Italian
r4	feta tomato olive oil	Greek
r5	yogurt cucumber garlic	Greek

- Estimate $P(\text{spaghetti} \mid \text{Italian})$, $P(\text{spaghetti} \mid \text{Greek})$, $P(\text{yogurt} \mid \text{Italian})$, and $P(\text{yogurt} \mid \text{Greek})$ according to the multinomial naive Bayes model for text classification. Use Laplace smoothing.
- Estimate the class prior probabilities $P(\text{Greek})$ and $P(\text{Italian})$.
- Use the probabilities estimated at (a) and (b) to compute $P(\text{Italian} \mid r6)$ according to the multinomial naive Bayes model, with r6: spaghetti yogurt.

Exercise 2: Naive Bayes for Text Classification

You are given the following collection of song lyrics with corresponding music genre:

songID	words in lyrics	Genre
s1	shake ya funky funky ya ya	Funk
s2	shake baby shake	Funk
s3	fire hell thunder hell	Metal
s4	blood hell venom	Metal
s5	hell ya burn	?

Here s1-s4 are the training examples, and s5 is a test example with unknown class label.

- (a) Estimate $P(\text{Funk} \mid s5)$ en $P(\text{Metal} \mid s5)$ according to the multinomial Naive Bayes model. Use Laplace smoothing.
- (b) Repeat (a), but this time don't use Laplace smoothing.

Exercise 3: Logistic Regression

We analyse data of professional darts games in order to predict the winner. Only games of type “best of x legs” have been taken into account. In the model, the probability that the player who begins, player a , wins against player b depends on the difference in average and checkout percentage between the two players, and a constant β_0 :

$$P(a \text{ wins against } b) = \frac{\exp(\beta_0 + \beta_1(\text{Av}_a - \text{Av}_b) + \beta_2(\text{Check}_a - \text{Check}_b))}{1 + \exp(\beta_0 + \beta_1(\text{Av}_a - \text{Av}_b) + \beta_2(\text{Check}_a - \text{Check}_b))}$$

Here Av_x denotes the average of player x , Check_x the checkout percentage of player x , and a denotes the player who starts throwing in the first leg. Estimation by maximum likelihood yields the following results:

Coefficient	Estimate
β_0 (Intercept)	0.120
β_1	0.135
β_2	0.025

Answer the following questions:

- (a) How big is the advantage of the right to start the game according to the fitted model?
- (b) Do the signs of the coefficient estimates make sense? Explain.
- (c) Michael van Gerwen has average respectively checkout percentage of 102.7 and 46.2%. Vincent van de Voort has average respectively checkout percentage of 92.6 and 40.4%. According to the model, what is the probability that Michael van Gerwen wins against Vincent van de Voort if van Gerwen starts?
- (d) What if van de Voort starts?
- (e) Give the linear classification rule corresponding to this model that predicts the player most likely to win.